

Who am I?

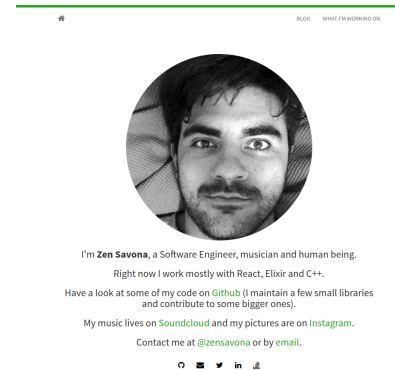
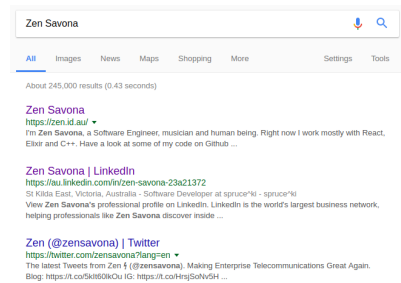
My name is Zen

I work for an enterprise telecommunications startup: spruce^ki

I think Elixir is pretty cool

@zensavona on Twitter

How do websites rank in Google? (simplified)



- Domains with more high quality links have more Link Juice™
- Link Juice™ flows from page to page
- More relevant link juice == higher ranking (sort of)
- Proprietary metrics to judge domains (TF, CF, DA, PA, etc)

The problem:

- There are lots of domains which expire and are full of Link Juice™.
- If we register these and link them to our site, we can hack ranking

but how can we find them?


The solution:

- Find a big website with a lot of Link Juice™
- Crawl it
- Find outward linked domains, throw away registered ones
- Throw away ones with shit metrics
- ?????
- Profit

Overview of architecture

- **Store:** Get stuff in and out of Redis
- **Scraper:** Core scraping and domain checking functionality
- **Workers:** Do the actual work and scale concurrency
- **Scheduler:** Dish each crawl an equal slice of the resource pie
- **Finisher:** Finish things up and put them in Postgres
- **Web:** Show me what's going on

Granular MicroService Driven Architectual Design Specification



Store

```
pool_size = 500
redix_workers = for i <- 0..(pool_size - 1) do
  worker(
    Redix,
    [
      [
        host: Application.get_env(:store, :redis_host),
        port: Application.get_env(:store, :redis_port)
      ],
      [name: :redix_#{i}]
    ],
    id: {Redix, i}
  )
end
```

```
def command(command) do
  Redix.command(:redix_#{random_index()},
               command, timeout: :infinity)
end
```

Scraper

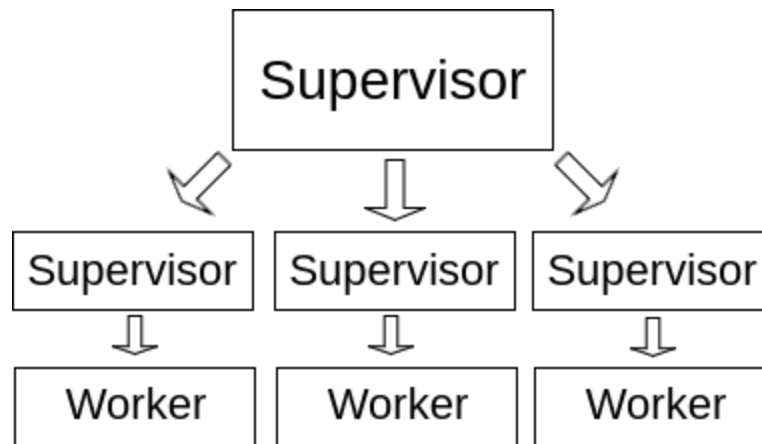
```
def check_domain(domain) do
  with {:ok, domain} <- domain_kind_of_at_least_makes_sense?(domain),
       {:ok, parsed} <- Domainatrex.parse(domain),
       domain <- "#{Map.get(parsed, :domain)}.#{Map.get(parsed, :tld)}",
       {:error, %HTTPoison.Error{id: nil, reason: :nxdomain}} <- HTTPoison.get(domain, [], hackney: [pool: :first_pool]),
       {:ok, %Whois.Record{created_at: nil}} <- Whois.lookup(domain),
       :available <- check_from_dnsimple(domain) do
    {:available, lookup_stats(domain)}
  else
    {:ok, _http} -> {:registered, nil}
    _ -> {:error, nil}
  end
end
```

```
def url_to_urls_and_domains(url) do
  domain = url |> domain_from_url

  with {:ok, %HTTPoison.Response{status_code: 200, body: body, headers: headers}} <- HTTPoison.get(url, [], hackney: [pool: :first_pool]),
       true <- html_content_type?(headers) do
    Floki.find(body, "a") |> Floki.attribute("href") |> normalise_urls("http://#{domain}") |> links_to_urls_and_domains(domain, url)
  else
    {:ok, %HTTPoison.Response{status_code: 301, headers: headers}} ->
      url = headers |> Enum.into(%{}) |> Map.get("Location")
      {:ok, [url], []}
    {:ok, %HTTPoison.Response{status_code: 302, headers: headers}} ->
      url = headers |> Enum.into(%{}) |> Map.get("Location")
      {:ok, [url], []}
    _ ->
      {:error, url}
  end
end
```

Workers

- One for one supervision, with a parent supervisor
- Simple, but works just fine



- Workers are just tail recursive Tasks, which ask the Store for a new thing when they finish working on the current thing

```
worker(Task,  
    [&Workers.Url.worker/0],  
    [id: {Workers.Url, id},  
     restart: :permanent]  
)
```


What do these workers do?

```
def worker do
  case Scheduler.pop_url do
    :empty ->
      :timer.sleep(1000)
    {crawl_id, url} ->
      case Scraper.Core.url_to_urls_and_domains(url) do
        {:error, url} ->
          Store.Crawled.push(crawl_id, url)
        {:ok, urls, domains} ->
          Store.Crawled.push(crawl_id, url)
          Store.ToCrawl.push(crawl_id, urls)
          domains
          |> Enum.each(
            &(Store.Domains.push(crawl_id, &1))
          )
      end
    end
  end
  worker()
end
```

Scheduler

Problems:

- If we pop a random url, bigger crawls would get more time
- If we pop the first/last, older crawls might never finish

Solution:

- Select a random crawl
- Select a random url from that crawl's `to_crawl` store

Finisher(s)

- How do we know when a crawl is finished?
- Domain metrics APIs cost \$500/mo (each), too expensive
- There is a "pirated" domain metrics API
- This API is really crap, needs retries
- The Finishers are just a `:timer`. Simple, but works fine

```
:timer.apply_interval(5000, Finisher, :finish, [])
```

Web

- Just a Phoenix app which adds crawls and seed urls to Redis

Crawls

Domains

Sign out

Crawl for <https://www.gwern.net/>

completed in 42 sec

6981
URLS CRAWLED

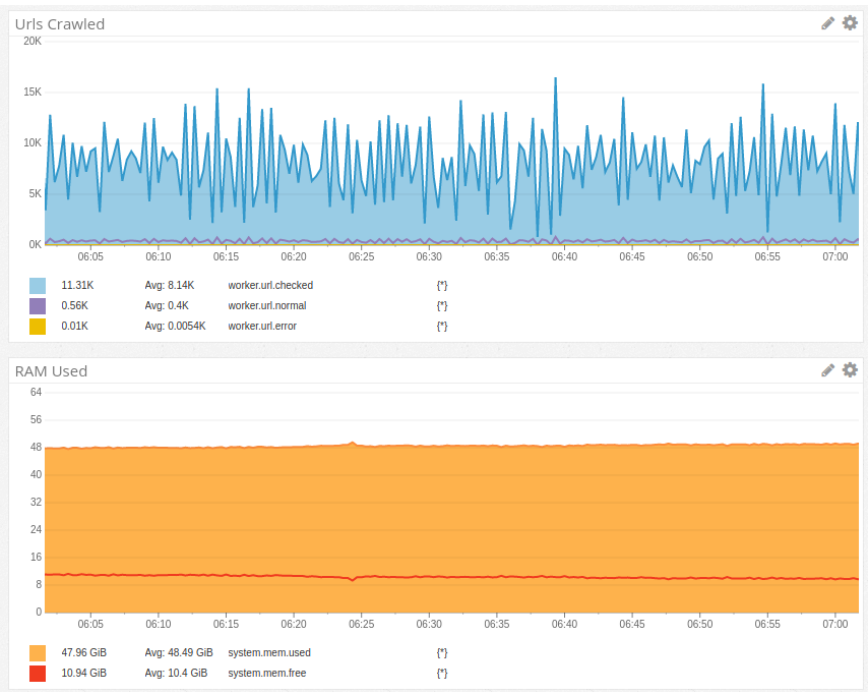
2069
DOMAINS FOUND

8
EXPIRED DOMAINS

Domain	Status	Domain Authority	Page Authority	Trust Flow	Citation Flow	MozRank
tvtropesf.org	Available	5.8	1.0	0.0	0.0	0.0
ganjisaffar.com	Available	6.0	1.0	0.0	6.0	0.0
holidaymead.com	Available	7.3	1.0	1.0	6.0	0.0
brandimontelab.it	Available	7.7	1.0	1.0	10.0	0.0
fuzziebutter.com	Available	9.9	1.0	4.0	14.0	0.0
thenationalbusinessassociation.com	Available	20.0	30.2	8.0	5.0	0.0
teambrainz.com	Available	7.5	1.0	0.0	0.0	0.0
theoverviewblog.com	Available	8.8	1.0	3.0	13.0	0.0

Instrumentation

- Datadog for graphs
- DogstatsD for metrics collection
- Looks pretty sweet



Problems / Solutions

- Redis is single threaded
- Redis is not as fast as you'd think with 100m+ Sets
- RAM is expensive, SSD is slow(?)

```
127.0.0.1:6379> slowlog get 5
1) 1) (integer) 318
   2) (integer) 1494995126
   3) (integer) 116336414
   4) 1) "flushall"
2) 1) (integer) 317
   2) (integer) 1494926004
   3) (integer) 19380
   4) 1) "SISMEMBER"
      2) "crawled_187"
      3) "https://www.meetup.com/fr-FR/topics/cultural-diversity-and-personal-growth/il/"
3) 1) (integer) 316
   2) (integer) 1494925592
   3) (integer) 13562
   4) 1) "SADD"
      2) "to_crawl:187"
      3) "https://www.meetup.com/fr-FR/Women-in-Software-Engineering-Los-Angeles/"
```

Questions / Comments / Improvements